

Notebook: CNNs in PyTorch

Now, you're ready to define and train an CNN in PyTorch.

To open this notebook, you have two options:

- Go to the next page in the classroom (recommended).
- Clone the repo from [Github](#) and open the notebook **cifar10_cnn_exercise.ipynb** in the **convolutional-neural-networks > cifar-cnn** folder. You can either download the repository with git clone <https://github.com/udacity/deep-learning-v2-pytorch.git>, or download it as an archive file from [this link](#).

Instructions

- Define a CNN model for classifying CIFAR10 images
- Train it for some number of epochs and test your model to see how well it generalizes and measure its accuracy.

This is a self-assessed lab. If you need any help or want to check your answers, feel free to check out the solutions notebook in the same folder, or by clicking [here](#).

GPU Workspaces

The next workspace is **GPU-enabled**, which means you can select to train on a GPU instance. The recommendation is this:

- Load in data, test functions and models (checking parameters and doing a short training loop) while in CPU (non-enabled) mode
- When you're ready to extensively train and test your model, **enable** GPU to quickly train the model!

All models, and the data they see as input, will have to be moved to the GPU device, so take note of the relevant movement code in the model creation and training process.